

SIMD in JavaScript via C++ and Emscripten

Peter Jensen

Intel Corporation
peter.jensen@intel.com

John McCutchan

Google Inc.
johnmccutchan@google.com

Ivan Jibaja

Intel Corporation
ivan.jibaja@intel.com

Dan Gohman

Mozilla
sunfish@mozilla.com

Ningxin Hu

Intel Corporation
ningxin.hu@intel.com

Abstract

This is the text of the abstract.

1. Introduction

We'll explore the use of Mozilla's Emscripten to compile C++ programs, that has use of SIMD intrinsics or gcc style vector code, into JavaScript. This was recently made possible by the SIMD.JS primitives introduced in JavaScript engine prototypes for Chromium and Firefox as well as extensions to the Emscripten compiler. Emscripten will correctly translate a subset of available C++ SIMD x86 intrinsics into corresponding operations defined in SIMD.JS. The JavaScript benchmarks associated with the SIMD.JS primitives were converted to C++ by hand, and then automatically converted back into JavaScript using the Emscripten compiler.

2. SIMD.JS

SIMD is short for Single Instruction, Multiple Data. It refers to CPU instruction level data parallelism. Most modern CPUs have a significant portion of their available instructions dedicated to operating on data in parallel. Typically, those instructions will perform the same operation on elements in short vectors, e.g. vectors of length 4, 8, or 16. Use of these instructions leads to increased performance, because more data processing is achieved with fewer instructions executed, and fewer instructions also means power savings, which is of outmost importance on mobile battery powered devices. Figure 1 shows how four scalar additions are combined into a single operation.

JavaScript is quickly emerging as one of the most popular languages among software developers. It was originally used for simple web page scripting for creating interactive web pages. Around 2008, very efficient and high performance JavaScript engines emerged, e.g. Firefox's TraceMonkey and Chrome's V8 engines. Since then, JavaScript has become a viable language for things beyond just basic web page interactivity, as witnessed by it's use in

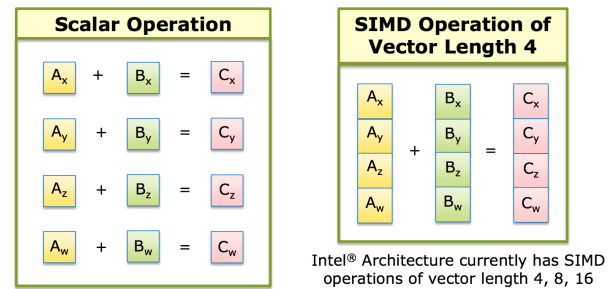


Figure 1. Replacing four scalar additions with one SIMD addition

large web based applications, such as office applications; e-mail, document processing, etc. Also, large games, which were previously standalone, natively compiled programs, have been ported to JavaScript to run within the browser environment. More recently, JavaScript has been adopted as a server side scripting language (node.js), and lately, JavaScript has found it's way to the mobile platform as a language that offers better portability between the different mobile platforms without sacrificing performance and features. For example, access to platform sensors (location, accelerometers, etc) are accessible from JavaScript via W3C APIs.

Even with the past 7 years of JavaScript performance advances, the desire for better performing JavaScript engines has not lessened, quite the contrary. It's a spiral that keeps on going; better performance leads to more uses, more uses require better performance. Specifically, software that use data parallelism to achieve adequate performance have, so far, been restricted to natively compiled languages, such as C++, because such languages offer ways of utilizing the SIMD instructions available in modern CPUs. JavaScript has only one number type, Number, which is an IEEE-754 floating point number, and JavaScript offers no abstraction primitives for writing algorithms utilizing data parallelism, so it's imperative that this shortcoming is dealt with, such that the next leap in JavaScript performance is made possible. This is what the SIMD.JS proposal addresses.

SIMD.JS is an emerging standard developed collaboratively by Intel, ARM, Mozilla, Google, and Microsoft. It provides low level data types and operations that map well onto the available SIMD instructions of the underlying hardware. Currently, the defined data types and operations are a representative and useful overlap between SIMD types and operations available in most modern CPUs.




Figure 2. SIMD.JS object hierarchy

```

1 function average(data) {
2   var sum = SIMD.float32x4.splat(0.0);
3   for (var i = 0, l = data.length; i < l; i = i+4) {
4     sum = SIMD.float32x4.add(
5       sum, SIMD.float32x4.load(data, i));
6   }
7   var total = sum.x + sum.y + sum.z + sum.w;
8   return total/data.length;
9 }

```

Figure 3. SIMD JavaScript code for finding the average of an array of numbers

The SIMD.JS proposal is structured as an object hierarchy, with SIMD being the top level global object. The immediate properties of the SIMD object reflect the data types; `int32x4`, `float32x4`, and `float64x2`. The operations are methods declared as properties on the data type properties as outlined in Figure 2, which shows a portion of the object hierarchy.

We’ve modelled the semantics of the SIMD types and operations as a polyfill [REF]. This allows programmers to experiment without using a JavaScript engine that natively supports SIMD.JS. The polyfill also serves as documentation for the semantics and interfaces. It will also reflect the current state of the proposal. The proposal is under active development and changes are likely to happen as the proposal is being refined and moves forward through the approval process.

As an example use case, Figure 3 shows the SIMD JavaScript code for computing the average of an array of floating point numbers. The numbers are held in a `Float32Array` typed array; `data`. The benefit of using SIMD operations, for computing the average, is that four numbers can be added in one operation, thereby reducing the number of iterations by a factor of 4, and achieving an equivalent speedup.

The optimizing Just-In-Time (JIT) compiler in our Chrome/V8 SIMD enabled prototype is able to produce the code in Figure 4 for the body of the loop. The code shows how the compiler is able to utilize 128-bit SIMD registers (`xmm`) to hold the value of `sum` and to use the `addps` instruction for adding 4 single precision numbers in one instruction. The details for this code snippet are as follows:

line 1-2: Check the loop index, `i`, against the upper bound and exit the loop if upper bound is encountered. `eax` holds the loop index and `edx` holds the upper bound.

line 3-4: Check to enable the JavaScript engine to abort the execution, if the loop has been running for too long. It will prevent a user program from hanging the browser.

line 5-11: Bounds check for the `SIMD.float32x4.load` operation. `eax` holds the index and `ecx` holds the upper bound.

line 12: Load the 4 32-bit float values. `edi` holds the base address of the data. The reason the index is multiplied by 2 and not

```

1 371 cmp eax,edx
2 373 jnl 440 (270629B8)
3 379 cmp esp,[0x645e9c]
4 385 jc 445 (270629BD)
5 391 mov esi,eax
6 393 shl esi,1
7 396 add esi,0x10
8 399 mov ebx,ecx
9 401 shl ebx,1
10 404 cmp esi,ebx
11 406 ja 499 (270629F3)
12 412 movups xmm2,[edi+eax*2]
13 416 movaps xmm3,xmm1
14 419 addps xmm3,xmm2
15 422 mov ebx,eax
16 424 add ebx,0x8
17 427 jo 504 (270629F8)
18 433 movaps xmm1,xmm3
19 436 mov eax,ebx
20 438 jmp 371 (27062973)

```

Figure 4. JIT compiler generated code for the for loop in the average function

4 is the fact that V8 represents integers in bits 1-31, so the value in `eax` is already holding the value of the index variable times 2.

line 13-14: Add the four 32-bit float values.

line 15-17: increment the loop index by 4. Since the integer representation used by V8 is done in bits 1-31 the actual value added is 8. The overflow checks is there to ensure that the result can still be represented in 31 bits, if not the representation is switched to a boxed floating point number.

line 18: Move the result of the SIMD add operation, to the `xmm1` register, which holds the value of `sum`.

line 19-20: Move the new loop index value to it’s proper register and jump back to the top of the loop.

For more details on how the JIT compilers operate see [REF].

2.1 The Future of SIMD.JS

The proposal has been presented to TC39, the JavaScript language standard committee, and was unanimously approved for stage 1 in 2014. Stage 1 is the proposal stage. It indicates that the need has been justified, and an outline for a solution has been accepted. It does not mean that this is the final proposal.

The focus, so far, has been on identifying types and operations that can be effectively implemented on all relevant CPU architectures. We realize that CPUs have distinct features that are useful and it will make sense to expose such features to the JavaScript programmer. This will most likely be done via architecture specific extensions to the SIMD object, e.g. `SIMD.x86`.*

SIMD.JS is currently being refined and prepared for the next stages of approval, and we expect this to be part of the EcmaScript 7 standard (ES7). EcmaScript 5 is the current JavaScript standard. EcmaScript 6 is slated for a mid-2015 release. ES6 is a major overhaul of the JavaScript language and a substantial set of new features were added, as reflected by the size of the language specification document. The ES5 specification document is roughly 300 pages, whereas the ES6 specification is roughly double that. Most browsers have already implemented most of the ES6 features.

3. Emscripten

Emscripten is a compiler that compiles C/C++ programs into JavaScript. It is based on the clang/LLVM compiler infrastructure [REF]. The compiler is the brainchild of Alon Zakai of Mozilla.

As an example of how it works, we’ll look at the generated JavaScript code resulting from compiling a simple C function.

```

1 float averageScalar(float *a, uint32_t length) {
2     float sum = 0.0f;
3     for (uint32_t j = 0, l = length; j < l; j = j + 4) {
4         sum = sum + (*(a++));
5     }
6     return sum/length;
7 }

```

Figure 5. Scalar C code for the average function

```

1 function _averageScalar($a, $length) {
2     $a = $a | 0;
3     $length = $length | 0;
4     var $a$addr$06 = 0, $add = 0.0, $j$05 = 0,
5         $sum$0$lcssa = 0.0, $sum$04 = 0.0, sp = 0;
6     sp = STACKTOP;
7     if (($length | 0) == 0)
8         $sum$0$lcssa = 0.0;
9     else {
10        $a$addr$06 = $a;
11        $j$05 = 0;
12        $sum$04 = 0.0;
13        while (1) {
14            $add = $sum$04 + +HEAPF32[$a$addr$06 >> 2];
15            $j$05 = $j$05 + 4 | 0;
16            if (($j$05 >>> 0 < $length >>> 0)) {
17                $sum$0$lcssa = $add;
18                break;
19            } else {
20                $a$addr$06 = $a$addr$06 + 4 | 0;
21                $sum$04 = $add;
22            }
23        }
24    }
25    STACKTOP = sp;
26    return +($sum$0$lcssa / +($length >>> 0));
27 }

```

Figure 6. JavaScript code generated by Emscripten for the averageScalar function

We'll again use a function that computes the average of float numbers. Figure 5 shows the input C program.

The Emscripten compiler command is similar to the clang compiler command, and takes most of the same options. The following command will generate optimized JavaScript code shown in Figure 6:

```
$ emcc -O2 -g average-scalar.c
```

This example shows how Emscripten manages to map a statically typed language (C) with pointers to a dynamically typed language without pointers (JavaScript).

Memory is modelled as overlayed typed arrays. In this example when the pointer `*a` is used to fetch from memory the corresponding JavaScript code is `+HEAPF32[aaddr$06 >> 2]` (line 14). `HEAPF32` is a global JavaScript typed array declared as follows:

```

var buffer = new ArrayBuffer(TOTAL_MEMORY);
HEAP8 = new Int8Array(buffer);
HEAP16 = new Int16Array(buffer);
HEAP32 = new Int32Array(buffer);
HEAPU8 = new Uint8Array(buffer);
HEAPU16 = new Uint16Array(buffer);
HEAPU32 = new Uint32Array(buffer);
HEAPF32 = new Float32Array(buffer);
HEAPF64 = new Float64Array(buffer);

```

All of these typed arrays are views on the same array buffer, so they all access the same physical memory. Notice that the index expression `'aaddr$06 >> 2'` is shifted right by 2. This is be-

```

1 float averageIntrin(float *a, uint32_t length) {
2     _m128 sumx4 = _mm_set_ps1(0.0);
3     for (uint32_t j = 0, l = length; j < l; j = j + 4) {
4         sumx4 = _mm_add_ps(sumx4, _mm_loadu_ps(a++));
5     }
6     float mSumx4[4];
7     _mm_storeu_ps(mSumx4, sumx4);
8     return (mSumx4[0] + mSumx4[1] +
9             mSumx4[2] + mSumx4[3])/length;
10 }

```

Figure 7. SIMD C code with intrinsics for the average function

cause `aaddr$06` is a byte address, and elements in the `HEAPF32` are 4 bytes each.

To enable the JavaScript JIT compilers to generate efficient code two type coercion tricks are used.

For integers and pointers the `'expr | 0'` is used to guarantee that the type of the resulting expression is a 32-bit integer. JavaScript semantics of the bitwise `|` expression dictate that the resulting expression is a 32-bit integer. A side effect of pointers being 32-bit integers is that compiled C/C++ programs are restricted to a 32-bit address space.

For floating point numbers, the unary `'+'` operator is applied, because JavaScript semantics dictate that the resulting expression is a floating point number.

Emscripten has been successfully used to compile very large C/C++ code bases (+100K lines of code). For example both Epic's and Unity's game engines have been ported, using Emscripten [REF]. Game engines are one example of software that will have optional implementations of performance critical portions of the code implemented using SIMD features. Since, JavaScript hasn't had a way of utilizing these powerful low level SIMD features of the CPU, Emscripten has not been able to compile these highly tuned implementations of the performance critical sections of the code. However, with the introduction of SIMD.JS, Emscripten will now be able to take full advantage of those. The next section covers how this is accomplished.

4. Compiling C++ with SIMD intrinsics

Figure 7 shows a typical SIMD implementation of the average function in C, using x86 SIMD intrinsics.

The `_m128` type holds 4 32-bit float numbers. The `_mm.*_ps` function calls are the SIMD intrinsics, which operates on single precision `_m128` values. For example, the `_mm_add_ps` intrinsic maps to the x86 `addps` instruction, which adds 4 32-bit float numbers in one operation. This allows the iteration count to be reduced by a factor of 4 resulting in an equivalent speedup.

The resulting JavaScript produced by Emscripten is shown in Figure 8.

The JavaScript code in Figure 8 might appear extensive at first look, however, it is very similar to the handwritten version of the function from Figure 3. The while loop starting at line 15 corresponds to the for loop from the C program. Implementing a for loop as a while loop takes a bit more code. The important thing to notice here is the use of the `SIMD_float32x4_add` and `SIMD_float32x4_load` SIMD.JS operations. The use of `'.'` instead of `'.'` is because Emscripten has created single identifier versions for all the SIMD.JS primitives.

It's important to note that use of CPU specific intrinsics makes the C version of the code target specific, i.e. it will only execute on x86 CPUs, whereas the resulting JavaScript code will execute on all architectures supported by the underlying SIMD enabled JavaScript engine.

```

1 function _averageIntrin($a, $length) {
2   $a = $a | 0;
3   $length = $length | 0;
4   var $add$i = SIMD_float32x4(0, 0, 0, 0),
5       $j$09 = 0,
6       $sumx4$0$lcssa = SIMD_float32x4(0, 0, 0, 0),
7       $sumx4$0i0 = SIMD_float32x4(0, 0, 0, 0),
8       sp = 0;
9   sp = STACKTOP;
10  if (($length | 0) == 0)
11    $sumx4$0$lcssa = SIMD_float32x4_splat(Math_fround(0));
12  else {
13    $j$09 = 0;
14    $sumx4$0i0 = SIMD_float32x4_splat(Math_fround(0));
15    while (1) {
16      $add$i =
17        SIMD_float32x4_add(
18          $sumx4$0i0,
19          SIMD_float32x4_load(buffer, $a + ($j$09 << 2) | 0))
20      $j$09 = $j$09 + 4 | 0;
21      if (($j$09 >>> 0 < $length >>> 0)) {
22        $sumx4$0$lcssa = $add$i;
23        break;
24      } else $sumx4$0i0 = $add$i;
25    }
26  }
27  STACKTOP = sp;
28  return +(($sumx4$0$lcssa.w + ($sumx4$0$lcssa.z +
29    ($sumx4$0$lcssa.x + $sumx4$0$lcssa.y))) /
30    +($length >>> 0));
31 }

```

Figure 8. Emscripten generated SIMD JavaScript code for the average function

```

1 typedef float floatx4 __attribute__((vector_size(16)));
2 float averageVectorSize(float *a, uint32_t length) {
3   floatx4 sumx4 = {0.0f, 0.0f, 0.0f, 0.0f};
4   floatx4 *ax4 = (floatx4 *)a;
5   for (uint32_t j = 0, l = length; j < l; j = j + 4) {
6     sumx4 = sumx4 + *(ax4++);
7   }
8   return (sumx4[0] + sumx4[1] +
9     sumx4[2] + sumx4[3])/length;
10 }

```

Figure 9. SIMD C code with `vector_size` for the average function

Use of non target specific SIMD code in C is possible via the gcc `vector_size` attribute. Emscripten also supports compiling such code. A non target specific version of the `average` function is shown in Figure 9. The generated JavaScript code is virtually identical to the code resulting from the C code using intrinsics. If possible, developers should be encouraged to write their SIMD code using this more universal syntax.

5. Benchmarks

We’ve created a set of benchmark kernels. The kernels are written in both C++ and JavaScript. Each kernel will have both a scalar implementation and a SIMD implementation.

The C++ SIMD implementation is done using x86 intrinsics. The C++ kernels are compiled with both the C++ clang compiler and with the Emscripten compiler. This gives us a basis for comparing SIMD/Scalar speedups for both C++ and JavaScript as well as C++/JavaScript performance differences. For the JavaScript execution we’ve used our two SIMD enabled JavaScript engines; V8 and SpiderMonkey.




Figure 10. SIMD vs. Scalar relative speedups

The benchmarks are written such that each kernel operation is executed as many times as it takes for it to run in about a second. This guarantees that the optimizing JavaScript JIT compilers kick in, which is essential for optimal performance.

The sources for the C++ benchmarks can be found here: [jREF_C](#), and the sources for the handwritten JavaScript benchmarks can be found here: [jREF_J](#)

6. Results

We’ve collected performance results for these combinations: C++, handwritten JavaScript, Emscripten produced JavaScript. For JavaScript execution we’re used both V8 and SpiderMonkey.

- Natively compiled C++
- Handwritten JavaScript executed with V8
- Handwritten JavaScript executed with SpiderMonkey
- Emscripten produced JavaScript from C++ implementation executed with V8
- Emscripten produced JavaScript from C++ implementation executed with SpiderMonkey

6.1 SIMD vs. Scalar

Figure 10 shows the relative SIMD vs. Scalar performance of the five combinations. Greater than 1 means that the kernel ran that the SIMD kernel ran that much faster than the corresponding scalar kernel.

6.2 Scalar C++ vs. JavaScript

Figure 11 shows the relative scalar C++ vs. JavaScript performance for each of the five combinations. Less than 1 means that the JavaScript kernel ran that much slower than the corresponding C++ kernel.

6.3 SIMD C++ vs. JavaScript

Figure 12 shows the relative SIMD C++ vs. JavaScript performance for each of the five combinations. Less than 1 means that the JavaScript kernel ran that much slower than the corresponding C++ kernel.




Figure 11. Scalar C++ vs. Javascript relative performance

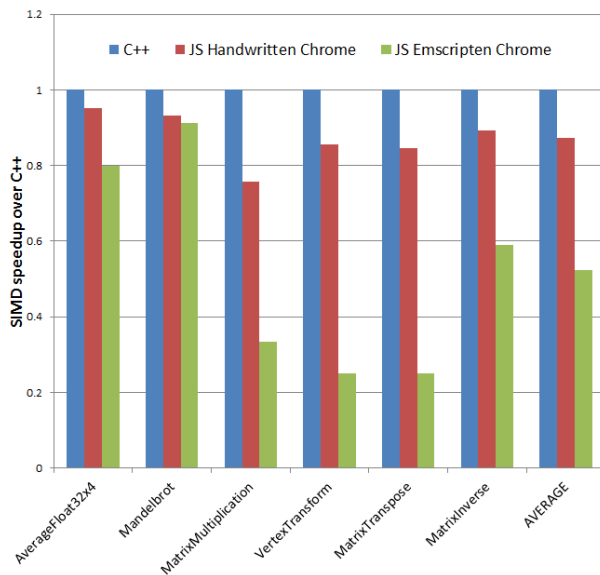


Figure 12. SIMD C++ vs. Javascript relative performance

7. Summary

A. Appendix Title

This is the text of the appendix, if you need one.

Acknowledgments

Acknowledgments, if needed.

References

[1] P. Q. Smith, and X. Y. Jones. ...reference text...